

The Config Assistant

Reimagining how customers configure AWS resources — from guessing in a blank form to being guided, before anything can break.

What I'm reimagining

Today, setting up an AWS resource means filling in fields on your own. The console rarely tells you if a value is dangerous until something has already gone wrong — and by then you're opening a support case. Help lives somewhere else: a docs page, a side panel, a forum you read afterward.

My reimagination puts a quiet assistant right inside the configuration screen. It reads every value as you set it, tells you plainly whether it's safe or risky, and suggests the right value when you're unsure — **before you save, not after you break something**.

Before and after

Today	With the Config Assistant
You fill in a blank form and hope the values are right.	Fields come pre-filled with safe defaults, each with a short reason.
You find out a setting was wrong when it breaks in production.	A risky value is flagged the moment you enter it, before you save.
Help is on a different page you have to go find.	The explanation is right beside the field you're editing.
Every customer relearns the same lessons the hard way.	It draws on both AWS-wide knowledge and your own account's history.

The core: a status box on every line

At the end of each configuration line sits a small box. Its icon reflects what the assistant thinks of your value, and clicking it explains why:

- **! (risky)** — the value is likely to cause a problem. Click it to see why, with a one-click button to change it to the safe value.
- **? (want a suggestion)** — you haven't set it yet, or you're unsure. Click it and the assistant recommends a value and explains its reasoning.
- **✓ (safe)** — the value looks good. Click it for a quick confirmation of why it's a sound choice.

Starting from scratch: the describe box

At the top of the screen is a box you can open and talk to. You describe what you're deploying in plain words — "a public web API," "an internal admin tool" — and the assistant recommends a full set of

safe defaults for your environment, which you can apply to the whole form in one click. It means a first-time user never has to face a blank, intimidating form.

Where its knowledge comes from: two layers

- **Common across AWS** — a shared baseline of the mistakes that hurt customers everywhere (open ports, too-low timeouts). This protects you from problems you've never even hit yet.
- **From your own account** — guardrails built from your own past incidents. This layer is specific to you, and it's the part no competitor can copy, because it's made from your history.

Proactive, without being annoying

The goal is to help before the error, not to nag. So the assistant is quiet by design:

- The most common help is silent — safe defaults are simply pre-filled, so no warning is ever needed.
- Reasons sit quietly beside each field. A real warning only appears when a value is genuinely risky.
- A 24-hour mute lets you silence a warning while you work, but it comes back on its own so you never accidentally turn off a real protection for good.
- Every time you dismiss a warning, that's a signal that trains the assistant to stop showing the noisy ones.

How it works underneath

Two parts. A **memory** stores what to watch for, what went wrong before, the safe value, and how badly it hurt. A **watcher** reads what you're doing, matches it against the memory, and surfaces the right status before you save. Clear-cut dangers use exact rules; the gray areas use fuzzy matching ("this looks like something that bit you before").

Building it for the demo

You can't add features to the real AWS console, so the prototype is a self-built mock of one configuration screen with the assistant inside it — exactly the medium-fidelity, clickable prototype the brief asks for. The logic is scripted around one believable scenario (an EKS deployment); if the team later wires in a real model, only the matching function changes.

Why it's a real shift

This isn't a chatbot bolted onto the console. It changes configuration from a solo guessing game into a guided, guarded conversation — and it turns every past mistake, yours and everyone's, into protection that shows up at the exact moment it matters. Support stops being the thing you call after a problem, and becomes the thing that quietly prevents it.